

# Knuth-Morris-Pratt algorithm (KMP)

Salvador Roura

The main use of the KMP algorithm is to efficiently search for all the instances of a word  $w$  into a text  $t$ .

The naive C++ code to solve this problem would be like this:

```
void find_matches(string w, string t) {  
    int m = w.size();  
    int n = t.size();  
    for (int i = 0; i ≤ n - m; ++i)  
        if (t.substr(i, m) == w) cout << i << endl;  
}
```

For instance, if  $w = \text{"aba"}$  and  $t = \text{"aabaacaabaa"}$ , this code prints 1 and 7.

In general (unless  $m = 0$  or  $m \simeq n$ ), the code above has cost  $\Theta(m \cdot n)$ . Similar codes have also cost  $\Theta(m \cdot n)$ , at least on the average.

We can do better. An alternative solution is the Rabin-Karp algorithm, which uses a rolling hash to solve this problem in  $\Theta(m + n)$  cost. (If you are curious, you can find the details online.)

However, the Rabin-Karp algorithm has the following faults (in comparison to KMP):

- The code is a bit error-prone.
- When the hash codes are equal, if we check if  $w$  trully matches the current position of  $t$ , the the cost can become  $\Theta(m \cdot n)$ . And if we don't perform this extra check, the algorithm has a small (but non-zero) probability of giving false positives.
- KMP is also a liner-time algorithm, but its constant is usually smaller, because it avoids expensive operations like modulus, which are typical of hash functions.

Before we present the KMP algorithm, we need a few definitions:

A *prefix* of a string  $s$  is a substring of  $s$  that occurs at the beginning of  $s$ . For instance, let  $s = \text{"aabaacaabaa"}$ . The prefixes of  $s$  are  $\lambda$  (the empty string), "a", "aa", "aab", ..., "aabaacaaba", and "aabaacaabaa" ( $s$  itself).

A *suffix* of a string  $s$  is a substring of  $s$  that occurs at the end of  $s$ . With the example above, the suffixes of  $s$  are  $\lambda$ , "a", "aa", "baa", ..., "abaacaabaa", and "aabaacaabaa".

A *border* of a string  $s$  is a substring  $b$  of  $s$  that is a prefix and also a suffix of  $s$ , and such that  $b \neq s$ . With the example above, the borders of  $s$  are  $\lambda$ , "a", "aa", and "aabaa".

There is a simple property that is key to KMP: *The border of a border is a border*. In the example, the borders of "aabaa" are indeed  $\lambda$ , "a", and "aa".

Now, we are ready to present the KMP algorithm. This is one of its possible implementations:

```
vector<int> kmp(string s) {
    int n = s.size ();
    vector<int> P(n);
    int j = -1;
    for (int i = 0; i < n; ++i) {
        while (j >= 0 and s[j] != s[i]) j = (j ? P[j-1] : -1);
        P[i] = ++j;
    }
    return P;
}
```

If we call this procedure with  $s = \text{"aabaacaabaa"}$ , the content of the resulting vector  $P$  is

	0	1	2	3	4	5	6	7	8	9	10
$s$	a	a	b	a	a	c	a	a	b	a	a
$P$	0	1	0	1	2	0	1	2	3	4	5

For every position  $i$  between 0 and  $n - 1$ ,  $P[i]$  contains the length of the longest border of  $s[0..i]$ . For instance,  $P[4] = 2$  corresponds to "aa", the longest border of  $s[0..4] = \text{"aabaa"}$ .

Another key to KMP is the following: *Iterating  $P$  provides all the borders of a string*. In the example,  $P[n - 1] = P[10] = 5$  provides "aabaa",  $P[5 - 1] = 2$  provides "aa",  $P[2 - 1] = 1$  provides "a", and  $P[1 - 1] = 0$  provides  $\lambda$ .

The algorithm fills the vector  $P$  from left to right, and computes each  $P[i]$  using the values of  $P[j - 1]$  for  $1 \leq j \leq i$ . To understand how it works, suppose that

we want to compute  $P$  for a string  $s'$  equal to  $s$ , but with a character appended to its right. Note that the first 11 positions of  $P$  are those already computed, and that we only need to calculate  $P[i]$  for  $i = 11$ . The value of  $j$  at that moment is 5, just stored at  $P[10]$ .

First, assume  $s'[11] = 'c'$ , that is,  $s' = "aabaacaabaac"$ . Because  $P[10] = 5$ , we already know that  $s'[0..4] = s'[6..10]$ . The code compares  $s'[5]$  against  $s'[11]$ . Since both are  $'c'$ , we can deduce  $s'[0..5] = s'[6..11]$  with only one character comparison, and so the longest border ending at  $s'[11]$  has length 6. Therefore, the **while** stops and we store 6 at  $P[11]$ .

Now, assume  $s'[11] = 'b'$ , that is,  $s' = "aabaacaabaab"$ . The comparison of  $s'[5]$  against  $s'[11]$  fails. Consequently, the code tries the next border, whose length can be found at  $P[j-1] = P[4] = 2$ . Now we successfully compare  $s'[2]$  against  $s'[11]$ , so  $P[11] = 3$ .

Similarly, when  $s'[11] = 'a'$  we get  $P[11] = 2$  after three iterations, and when  $s'[11] = 'd'$  we get  $P[11] = 0$  after four iterations. Note that this last case requires some care to end the **while** avoiding an access to  $P[-1]$ .

What is the cost of the KMP algorithm? An elemental reasoning provides the upper bound  $O(n^2)$ : we have  $n$  iterations, each with cost  $O(n)$ . Let us do a tighter analysis. First, note that the cost is dominated by the number of times that the  $j = (j ? P[j-1] : -1)$ ; instruction is executed. Second, every time it is executed,  $j$  is decreased by at least 1. But how many times can we decrease  $j$ ? At most as many times as we increase it ( $++j$ ), that is, at most  $n$  times. Therefore, the cost is no more (nor less) than  $\Theta(n)$ .

We have yet to show how to use KMP to search for all the instances of a word  $w$  into a text  $t$ . This is how: For simplicity, assume that there is a special character that does not appear in  $w$  nor in  $t$ , say  $'\#'$ . First, we concatenate  $w$ , the special character and  $t$ . Afterwards, we call KMP over this string, and search for the length of  $w$  into the vector  $P$ :

```
void find_matches(string w, string t) {
    string s = w + "#" + t;
    vector<int> P = kmp(s);
    int m = w.size ();
    for (int i = m + 1; i < s.size (); ++i)
        if (P[i] == m) cout << i - 2*m << endl;
}
```

For instance, with the first example  $w = "aba"$  and  $t = "aabaacaabaa"$ , we build  $s = "aba#aabaacaabaa"$ . The resulting vector  $P$  will only have a 3 at the positions 7 and 13, corresponding to the two matches. We subtract  $m$  twice from  $i$  to make this algorithm functionally equivalent to the naive code, although much faster in general (linear versus quadratic).